

```
In [1]: import numpy as np
import pandas as pd
from scipy.stats import mode
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_curve, roc_auc_score


%matplotlib inline
```

```
In [2]: data = pd.read_csv('Churn_Modelling.csv').dropna(axis=1)
data.drop(columns=['Surname', 'RowNumber', 'CustomerId'], inplace=True)
data = pd.get_dummies(data, columns=['Geography', 'Gender'])
data.info()

churn_count = data['Exited'].value_counts()
temp_df = pd.DataFrame({
    'Exited': churn_count.index,
    'Counts': churn_count.values
})

plt.figure(figsize=(10, 6))
sns.barplot(x='Exited', y='Counts', data=temp_df)
plt.xticks(rotation=90)
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   CreditScore            10000 non-null  int64
1   Age                    10000 non-null  int64
2   Tenure                  10000 non-null  int64
3   Balance                 10000 non-null  float64
4   NumOfProducts          10000 non-null  int64
5   HasCrCard               10000 non-null  int64
6   IsActiveMember         10000 non-null  int64
7   EstimatedSalary        10000 non-null  float64
8   Exited                  10000 non-null  int64
9   Geography_France       10000 non-null  bool
10  Geography_Germany      10000 non-null  bool
11  Geography_Spain        10000 non-null  bool
12  Gender_Female          10000 non-null  bool
13  Gender_Male            10000 non-null  bool
dtypes: bool(5), float64(2), int64(7)
memory usage: 752.1 KB
```



```
In [3]: X = data.iloc[:, :-1]
        y = data.iloc[:, -1]

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
        print(X_train.shape)
        print(X_test.shape)
```

```
(8000, 13)
```

```
(2000, 13)
```

```
In [4]: X_test, X_val, y_test, y_val = train_test_split(X_test, y_test, test_size=0.5, random_state=42)
```

```
print(X_test.shape)
print(X_val.shape)
```

```
(1000, 13)
```

```
(1000, 13)
```

```
In [5]: def cv_scoring(estimator, X, y):
        return accuracy_score(y, estimator.predict(X))

models = {
    "SVC":SVC(),
    "Gaussian NB":GaussianNB(),
    "Random Forest":RandomForestClassifier(random_state=18)
}

for model_name in models:
    model = models[model_name]
    scores = cross_val_score(model, X, y, cv = 10,
                              n_jobs = -1,
                              scoring = cv_scoring)

    print("=="*30)
    print(model_name)
    print(f"Scores: {scores}")
    print(f"Mean Score: {np.mean(scores)}")
```

```
=====
SVC
Scores: [0.545 0.545 0.545 0.546 0.546 0.546 0.546 0.546 0.546 0.546]
Mean Score: 0.5457000000000001
=====
Gaussian NB
Scores: [0.619 0.591 0.619 0.603 0.639 0.613 0.614 0.616 0.607 0.592]
Mean Score: 0.6113
=====
Random Forest
Scores: [1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
Mean Score: 1.0
```


```
In [6]: rf_model = RandomForestClassifier(random_state=18)
        rf_model.fit(X_train, y_train)
        preds = rf_model.predict(X_val)
        print(f"Accuracy on train data by Random Forest Classifier: {accuracy_score(y_train, rf_model.predict(X_train))*100}")
```

```
print(f"Accuracy on test data by Random Forest Classifier: {accuracy_score(y_val, preds)*100}")

cf_matrix = confusion_matrix(y_val, preds)
plt.figure(figsize=(12,8))
sns.heatmap(cf_matrix, annot=True)
plt.title("Confusion Matrix for Random Forest Classifier on Test Data")
plt.show()
```

Accuracy on train data by Random Forest Classifier: 100.0

Accuracy on test data by Random Forest Classifier: 100.0



```
In [7]: y_pred = rf_model.predict(X_test)
print(f"Accuracy of the model on test dataset: {accuracy_score(y_test, y_pred)}")
```

```

cm = confusion_matrix(y_test, y_pred)
print(f'Confusion Matrix:\n{cm}')

report = classification_report(y_test, y_pred)
print(f'Classification Report:\n{report}')

y_proba = rf_model.predict_proba(X_test)[:, 1]
auc = roc_auc_score(y_test, y_proba)
fpr, tpr, _ = roc_curve(y_test, y_proba)

import matplotlib.pyplot as plt
plt.figure()
plt.plot(fpr, tpr, label=f'ROC curve (area = {auc:.2f})')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc='lower right')
plt.show()

```


Accuracy of the model on test dataset: 1.0

Confusion Matrix:

```
[[463  0]
 [ 0 537]]
```

Classification Report:

	precision	recall	f1-score	support
False	1.00	1.00	1.00	463
True	1.00	1.00	1.00	537
accuracy			1.00	1000
macro avg	1.00	1.00	1.00	1000
weighted avg	1.00	1.00	1.00	1000



In [ ]: